

The `alethe-dev` branch of `lean-smt`: a summary of the changes by code structure

September 2026

Abstract

The branch `alethe-dev` adds an Alethe proof checker to `lean-smt`, in 104 commits over the merge base `76b4eea` of `main`. This note describes what was added, organized by code structure rather than by chronology: the parser, the changes to the pre-existing library, the new verified modules placed inside that library, and the reconstruction of the Alethe rules. For each part it gives its size, and for the rule reconstruction it lists, rule by rule, which existing `lean-smt` theorem, tactic, verified program or meta-level reconstructor is leveraged.

1 Overview

The additions fall into five parts. Three are the ones one expects, namely a parser for Alethe, the reconstruction of Alethe rules into Lean terms, and changes to routines that `lean-smt` already had. Two more are worth separating out: a set of new *verified* modules that were placed inside the existing reconstruction tree because both the `cvc5` path and the Alethe path use them, and the driver that checks a proof step by step in the kernel, which has no counterpart in the tactic-based `cvc5` path. Figure 1 shows how a proof flows through these parts, and Table 1 gives their sizes. The pre-existing lemma libraries, that is every `Lemmas.lean`, every `Rewrites.lean` and `Prop/Core.lean` under `Smt/Reconstruct/`, are byte-identical to `main`.

2 The parser

The parser is self-contained and depends on nothing in the reconstruction library beyond the S-expression reader `Smt/Data/Sexp.lean`, which existed on `main`.

`Syntax.lean` (205 lines)

The abstract syntax of Alethe proofs, parametric in the representation τ of terms: step arguments `Arg`, `StepData` (id, rule, clause, premises, arguments), `Command` (assume, step, anchor, subproof) and `Proof`. It also defines the hash-consed `Arena` of S-expression `Nodes` over `Nat` ids, with the queries the realizer needs (whether a node is open under a binder, its size, its binder variables).

`Parser.lean` (508 lines)

Turns the S-expressions of a proof into a `Proof Nat` over the arena. Every binding construct is eliminated here through one environment: parallel `let`, `:named` (whose atoms are computed lazily and once, so that sharing survives), user-level `define-fun`, and the variable renamings and assignments of `anchor`. Binders close their variables, so a closed named term is resolved once for the whole proof.

`Realize.lean` (233 lines)

Turns the `Proof Nat` into a `Proof cvc5.Term`. The SMT-LIB problem and every term of the proof are parsed by `cvc5`'s own parser, through `lean-cvc5`'s `InputParser`. Shared closed subterms of the arena are sent once, bound with `:named`, and referred to by name




Figure 1: A proof through the checker. The upper row is new; the lower row is the pre-existing reconstruction library, which the checker reaches through `cvc5`'s term representation. The realizer hands every term to `cvc5`'s own parser precisely so that the term reconstructors apply unchanged.

afterwards; open subterms are printed inline. The point of this design is that the reconstructor never sees an Alethe-specific term: `Step.lits` and `Premise.lits` are arrays of `cvc5.Term`, so the whole `@[smt_term_reconstruct]` and `@[smt_sort_reconstruct]` machinery of `lean-smt` applies unchanged.

3 Changes to the pre-existing library

Eleven files that existed on `main` under `Smt/Reconstruct/` were modified, for a net of +1361 / -791 lines. Not all of it is optimization. Each change is classified below as a performance optimization of an existing `cvc5` routine (P), a generalization so that the Alethe checker can call an existing routine (G), a bug fix found through the Alethe evaluation (F), or a new definition (N). Figure 2 shows where the changes land; the classes are detailed in the following subsections and summarized in Table 2.

3.1 Performance optimizations (P)

These change the proof terms the `cvc5` path produces, never its conclusions; each keeps the previous construction as a fallback where one is needed.

- **The polynomial normalizer sorts once.** In `Int/Polynorm.lean`, `Rat/Polynorm.lean` and `Real/Polynorm.lean`, `Expr.toPolynomial` used to fold a sum through a sorted insertion, so that a sum of n terms cost the kernel $O(n^3)$ steps. It now collects the monomials in one pass (`Expr.flatten`, threading a sign for `Int` and a rational multiplier for `Rat` and `Real` so that negation, subtraction and division by a constant cost nothing) and sorts them with a bottom-up merge sort (`mergePairs`, `mergeAll`, `normalize`), all structurally recursive because the kernel does not unfold well-founded recursion. The definitions and their denotation lemmas were rewritten and reproved; the exported soundness theorem `denote_eq_from_toPolynomial_eq` is unchanged in statement and proof, and nothing was sorried. Commit 7258b45. Measured on a 400-variable sum: 25s and 8.9 GB became 1.3s.
- **polyNorm decides the equality.** The kernel tactic closed its goal with `Eq.refl` between two normalizer applications, which the kernel unifies by unfolding both sides in lockstep instead of evaluating each; on large coefficients that ran for over half an hour. It now applies `of_decide_eq_true` to `Eq.refl true`, as `nativePolyNorm` already did. Meta code only; `nativePolyNorm` itself is untouched. Commit 711e8e7.
- **Projection instead of list indexing.** In `Prop.lean`, the `cvc5` rules `AND_ELIM`,

Part	Files	Lines
Alethe parser	Syntax, Parser, Realize (3 new)	946
Rule reconstruction	Basic, Clause, Prop, Arith, UF, Quant, Polyeq, Lemmas, Rare, Rare/Basic (10 new)	3 160
Generated RARE table	Rare/Rules.lean (1 new; generator gen.py 445, input rewrites.eo 513)	1 380
Driver, frontend, tactic	Reconstruct, Frontend, Command, Tactic, Alethe.lean (5 new)	1 381
New verified modules in the existing tree	Prop/AciNorm 461, Prop/Proj 239, Prop/ResNorm 397 (3 new)	1 097
Changes to pre-existing library files	11 modified	+1 361 / -791
Registration	Smt.lean, Attribute.lean, Options.lean (3 modified)	+8
Tests	Test/Alethe: 48 test files over 94 problem and proof files	256 + 15 037
Unit tests	Test/Unit/AciNorm, Test/Unit/ResNorm (2 new)	126
Scripts	check-alethe.sh, setup-alethe.sh, rule-boxplots.py (3 new)	636
Report and notes	docs-alethe/: T _E X 1 246, Markdown 606, plus figures and CSVs	

Table 1: Size of each part of the branch. Lines are lines of Lean unless stated otherwise. The 104 commits are on top of 76b4eea.

NOT_OR_ELIM, CNF_AND_POS, CNF_AND_NEG, CNF_OR_POS and CNF_OR_NEG projected the i -th conjunct through `andN` and `List.getElem` with a decided bound, charging every step for the whole conjunction. They now walk the right-nested `And` the hypothesis already is, through the new `Prop/Proj.lean` (Section 4). `CNF_OR_POS` collapses to the single node `notOrSelf`. Commit 008ac35.

- **ACI_NORM through the reflective normalizer.** In `Builtin.lean`, an `and/or` layer is normalized by the verified `AciNorm` normalizer, evaluated by the kernel; other operators and any failure fall back to the previous AC rewriting tactic. Kernel time for a 200-atom layer drops from seconds to about 30 ms. Commit 28bb705.
- **Decidable instances and closed-term caching.** In `Smt/Reconstruct.lean`, `synthDecidableInstance` guards the instance search (whose own “failed to synthesize” used to escape and take down whole runs on QF_UFIDL) and, when the search fails, builds the instance compositionally through $\neg$, $\wedge$, $\vee$ and $\leftrightarrow$ so that it matches the instance a theorem over `[Decidable c]` elaborates to. The reconstruction state gains a `closedTermCache`, valid across binder scopes, enabled only when a driver sets `baseLCtx`; the `cvc5` path leaves it off.

3.2 Generalizations for sharing (G)

Pure refactors with no behavior change on the `cvc5` path, almost all in the first commit c69e4d5.

- A record `RewriteStep` (rule, arguments, list arguments, result, lazily reconstructed premises) in `Smt/Reconstruct.lean` decouples a rewrite step from `cvc5.Proof`. Every per-theory `reconstructRewrite` in `Builtin`, `Int`, `Prop`, `Quant`, `Rat` and `UF` was retyped to take it, and each `cvc5` call site wraps its proof with `RewriteStep.ofProof`. This is what lets an Alethe `rare_rewrite` or `qnt_join` step, which has no proof object, reach the existing rewrite reconstructors.
- `Prop.getResolutionResult` and `Prop.reconstructResolution` take the pivot polarity as




Figure 2: Where the branch touches the existing reconstruction tree. Every per-theory file received the same generalization of its rewrite reconstructor; the optimizations concentrate in the polynomial normalizers and in `Prop.lean`, the latter through the new `Proj` module. The lemma libraries were not modified.

a `Bool` rather than a `cvc5` term; `Prop.isNotOf` became public.

- In `Int.lean` and `Rat.lean`, `reconstructArithPolyNormRel` takes the premise's term and proof instead of a proof object, and the term-level half of the mult-sign reconstruction was factored out as `reconstructMulSignTerms`, the old name staying as a two-line wrapper (commit eab472d).
- `addTrustWith` lets a step be trusted with a custom trace message.

3.3 Bug fixes (F)

- `Quant.lean`: `QUANT_UNUSED_VARS` was wrong when a prefix binds the same variable twice or reorders (`cvc5` reuses the variable node for a name, and the body refers to the innermost binder), and again when every binder is dropped and the body is itself a quantifier. Each side is now instantiated from the binders of the other. Commits 3ca2f21, 4d88b28. A regression test with a repeated binder was added.
- `Int.lean`, `Rat.lean`: a relation with one `Int` side and one `Real` side was claimed by the `Int` handler and reconstructed ill-sorted. The `Int` handler now declines unless both sides are integer, and the `Rat` handler casts through a new `ratArg` helper.
- `UF/Congruence.lean`: the left- and right-associative congruence helpers replaced panicking `Option.get!` and `hs[0]!` with errors, so a premise that is not an equality is reported instead of becoming a dummy term the kernel then rejects. Commit ad066ae.

File	+	−	Classes
Smt/Reconstruct.lean	97	10	P G
Builtin.lean	50	40	G P
Int.lean	114	97	G F N
Int/Polynorm.lean	222	92	P
Prop.lean	134	142	G P
Quant.lean	67	54	G F
Rat.lean	121	86	G F N
Rat/Polynorm.lean	256	120	P
Real/Polynorm.lean	263	121	P
UF.lean	21	21	G
UF/Congruence.lean	16	8	F

Table 2: The modified pre-existing files under `Smt/Reconstruct/` and the classes of change each received.

3.4 New definitions (N)

`Int.intValue` and `Rat.constValue` read a coefficient written as Alethe writes it, as a negated literal (`- 1`) or a quotient of literals (`/ 1 4`); for `cvc5` input they agree with the previous direct reads. No new theorem was added to any pre-existing file.

4 New verified modules in the existing tree

Three modules were added under `Smt/Reconstruct/Prop/` rather than under `Smt/Alethe/`, because they are reflective checkers in the style of the polynomial normalizer and serve both paths. Each has a verified core of ordinary definitions and theorems, and a small meta layer that reifies an expression and applies the soundness theorem as a term.

Prop/Proj.lean (239 lines, 8 theorems)

Projection out of a right-nested $\wedge$ or $\neg\vee$ chain, and the builders of the CNF axioms (`mkAndProj`, `mkNotOrProj`, `mkOrInj`, `mkCnfAndPos`, `mkCnfAndNeg`, `mkCnfOrNeg`, `notOrSelf`). A `ProjChain` keeps the partially unrolled spine so that a step projecting many indices out of one hypothesis builds $O(n)$ nodes in total. Used by the `cvc5` rules named above and by the Alethe rules `and`, `not_or`, `and_pos`, `and_neg`, `or_pos`, `or_neg`, `onpoint` and `distinct_elim`.

Prop/AciNorm.lean (461 lines, 20 theorems)

One layer of $\wedge$ or $\vee$ is reified into a tree over indexed atoms and normalized to a bitset, `none` when the annihilator occurs; a layer of `xor` is normalized to the parity of its atoms. The kernel evaluates the normal form with its accelerated `Nat` bit operations. Used by `cvc5`'s `ACI_NORM` and by Alethe's `semilattice_simp` and `boolean_group_simp`.

Prop/ResNorm.lean (397 lines, 18 theorems)

A checker for resolution chains and the containment rules. A literal is a `Nat`, a clause is both a `List Nat` whose denotation unfolds definitionally to the stated proposition and a bitset on which each link of a chain is a few accelerated operations; the pivots are hints, so a wrong one fails only at the final containment. It depends on `propext` and `Quot.sound` only. Used by the Alethe clausal rules behind the option `smt.alethe.reflect`, default off; the `cvc5` path does not use it yet.

5 Reconstruction of the Alethe rules

The rule handlers are registered with a new attribute `alethe_rule_reconstruct`, added to the existing attribute framework in three lines of `Smt/Attribute.lean` and collected by the existing




Figure 3: What each rule family leverages. Grey boxes are components that existed on `main` and were not modified; orange ones existed and were generalized or factored out on the branch; green ones are the new verified modules of Section 4.

`getReconstructors`. Five handlers cover the rules, one per family (`Clause`, `Prop`, `Arith`, `UF`, `Quant`), with `Rare` dispatched from `UF`.

This section covers the rules that reach the checker after Carcara’s pipeline as the evaluation runs it: `polyeq`, `local`, `core-simp-rare` restricted by `-core-rules` to the `*_simplify` family and the legacy `ac_simp`, `aci_simp` and `absorb`, `budget`, and the `reordering` pass. Those passes remove the `*_simplify` rules and the legacy ACI rules (replaced by `rare_rewrite` and `evaluate` chains and by the structural `semilattice_simp`, `boolean_group_simp` and `assoc_simp`) and `reordering`. The non-linear multiplication rules are left out as well. The handlers also accept those rules, so that a raw solver proof can be checked, but they are not covered here. The reconstructors of `bfun_elim`, `th_resolution` and `strict_resolution` have been dropped, no corpus of the evaluation emitting them after the pipeline. Figure 3 shows which existing component each family draws on; the subsections list every rule. Throughout, a component is *existing* if it is present on `main` and *new* otherwise; the new verified modules of Section 4 are marked as such.

5.1 Substrate shared by every rule

Before any rule, the checker reuses the reconstruction monad and its services unchanged: `ReconstructM` with `Reconstruct.Context` and `Reconstruct.State`; the term and proof caches, `withNewTermCache` and `withNewProofCache`; the assumption handling, `withAssums` and `findAssumWithType?`; the step recorders `addThm`, `addTac` and `addTrust`, and the option `useNative`. Terms and sorts go through `reconstructTerm`, `reconstructSort` and `reconstructSortLevelAndSort`, hence through every `@[smt_term_reconstruct]` and `@[smt_sort_reconstruct]` handler of `Prop`, `Builtin`, `UF`, `Int`, `Rat`, `Real`, `BitVec` and `Datatype`. The `alethe` tactic reuses the entire `smt` front end: `Preprocess.mono`, `pushHintsToCtx`, `intros`,

negateGoal, normalize, embedding, applySteps, genUniqueFVarNames, prepareSmtQuery, Command.cmdsAsQuery and the Config/Result types. The problem's symbols are introduced with reconstructSort through a flat withLocalDecls rather than one nested withLocalDecl per declaration, and choice symbols are bound to Classical.epsilon over reconstructSortLevelAndSort.

5.2 Clausal rules (Clause.lean, 227 lines)

Rule	Component
resolution	existing Prop.reconstructResolution, Prop.getResolutionResult; new ResNorm.proveClause when smt.alethe.reflect is set
contraction, or weakening	new lemma orN_of_map; ResNorm when set existing Prop.orN_append_left
subproof	new lemma orN_of_impliesN
hole	trusted, through addTrustWith

The fix-up that reconciles a computed clause with the stated one (permutations, duplicates, flattening) uses the existing Prop.eqResolve and the existing AC tactic Meta.AC.rewriteUnnormalizedTop from Builtin/AC.lean as its last resort. Resolution on a premise that itself holds the pivot no longer re-resolves. The existing resolution reconstructor is the largest single reuse in the checker.

5.3 Propositional rules (Prop.lean, 266 lines)

Of the 45 rules, 40 are direct applications of lemmas that existed in Prop/Lemmas.lean and Builtin/Lemmas.lean.

Rule	Component
not_and	existing Prop.notAnd
implies	existing Prop.impliesElim
not_implies1, not_implies2	existing Prop.notImplies1, Prop.notImplies2
equiv1, equiv2	existing Prop.equivElim1, Prop.equivElim2
not_equiv1, not_equiv2	existing Prop.notEquivElim1, Prop.notEquivElim2
xor1, xor2	existing Prop.xorElim1, Prop.xorElim2
not_xor1, not_xor2	existing Prop.notXorElim1, Prop.notXorElim2
ite1, ite2	existing Builtin.iteElim2, Builtin.iteElim1
not_ite1, not_ite2	existing Builtin.notIteElim2, Builtin.notIteElim1
implies_pos	existing Prop.cnfImpliesPos
implies_neg1, implies_neg2	existing Prop.cnfImpliesNeg1, Prop.cnfImpliesNeg2
equiv_pos1, equiv_pos2	existing Prop.cnfEquivPos2, Prop.cnfEquivPos1
equiv_neg1, equiv_neg2	existing Prop.cnfEquivNeg2, Prop.cnfEquivNeg1
xor_pos1, xor_pos2	existing Prop.cnfXorPos1, Prop.cnfXorPos2
xor_neg1, xor_neg2	existing Prop.cnfXorNeg2, Prop.cnfXorNeg1
ite_pos1, ite_pos2	existing Builtin.cnfItePos2, Builtin.cnfItePos1
ite_neg1, ite_neg2	existing Builtin.cnfIteNeg2, Builtin.cnfIteNeg1
and, not_or	new Proj.mkAndProj, Proj.mkNotOrProj
and_pos, and_neg, or_neg	new Proj.mkCnfAndPos, mkCnfAndNeg, mkCnfOrNeg
or_pos	new Proj.notOrSelf
true	trivial
false, not_not	new lemmas not_false_cl, not_not_cl
ite_then_intro, ite_else_intro	new lemmas of the same names
connective_def, qnt_duality	new lemmas connective_def_xor/eq/ite; quantifier cases by a small simp set; existing synthDecidableInstance
and_intro	new meta code, an And.intro fold

5.4 Arithmetic rules (Arith.lean, 553 lines)

The Farkas driver behind la_generic, la_tautology and bounded_farkas is new meta code, about 340 lines: the negations of the clause's literals are assumed, oriented into bounds, strength-

ened for strict integer bounds, scaled by the coefficients given as `:args`, summed, and refuted through the normalized difference of the two sides and a decidable fact about the resulting constant. Every proof-level ingredient of that driver is existing.

Rule	Component
<code>la_generic</code> , <code>la_tautology</code> , <code>bounded_farkas</code>	existing <code>Int.sum_ub1..9</code> , <code>Rat.sum_ub1..9</code> ; existing <code>mul_pos_lt/le/eq</code> , <code>mul_neg_eq</code> (<code>Int</code> , <code>Rat</code>); existing <code>Int.polyNorm</code> , <code>Rat.polyNorm</code> and their native variants; existing <code>Prop.notNotElim</code> ; Lean core <code>omega</code> for integer tightening; new lemmas <code>farkas_lt/le/eq</code> (<code>Int</code> , <code>Rat</code>), <code>orN_of_impliesN_not</code>
<code>poly_simp</code> <code>poly_simp_rel</code>	existing <code>Int.polyNorm</code> , <code>Rat.polyNorm</code> , and native variants <code>reconstructArithPolyNormRel</code> (<code>Int</code> , <code>Rat</code>), generalized from the existing handler; nine new mixed-sort lemmas over the existing <code>Rat.*_of_sub_eq*_int_left</code>
<code>la_rw_eq</code>	new lemmas <code>Int.la_rw_eq</code> , <code>Rat.la_rw_eq</code> over existing <code>Rat.trichotomy₂</code>
<code>la_disequality</code>	new lemmas over existing <code>Rat.trichotomy₂</code>
<code>la_totality</code>	new lemmas <code>Int.la_totality</code> , <code>Rat.la_totality</code>
<code>div_intro</code>	new lemmas <code>div_intro_pos/neg/real</code> ; <code>omega</code> for a constant divisor
<code>div_by_zero_intro</code>	new lemma <code>div_by_zero_intro</code>
<code>lia_generic</code>	Lean core <code>omega</code> only; a step outside its fragment is trusted

The new Farkas, disequality, totality and `la_rw_eq` lemmas in `Alethe/Lemmas.lean` are themselves proved from existing `Rat.trichotomy2`, `Rat.lt_eq_sub_lt_zero`, `Rat.le_eq_sub_le_zero`, `Rat.lt_of_le_of_lt`, `Rat.lt_of_lt_of_le` and `Rat.le_of_lt`. No lean-smt tactic is invoked by name: the polynomial normalizer is called as a function on a goal, and `omega` is Lean's.

5.5 Equality and rewriting rules (`UF.lean`, 556 lines)

Rule	Component
<code>cong</code> <code>eq_congruent</code> , <code>eq_congruent_pred</code>	existing <code>UF.smtCongr</code> , <code>UF.eq_symm</code> ; new <code>transportEq</code> for distinct existing <code>UF.smtCongr</code> , <code>Prop.notNotElim</code> ; new lemma <code>orN_of_impliesN_not</code>
<code>eq_transitive</code> <code>eq_symmetric</code>	existing <code>Prop.notNotElim</code> ; new chaining over <code>Eq.trans</code> existing <code>UF.eq_symm</code> , <code>Prop.impliesElim</code>
<code>refl</code> , <code>eq_reflexive</code> <code>symm</code> , <code>not_symm</code> , <code>trans</code>	existing <code>reconstructSortLevelAndSort</code> new meta code over <code>Eq.symm</code> , <code>Ne.symm</code> , <code>Eq.trans</code> , with defeq-aware orientation under anchor <code>lets</code>
<code>distinct_elim</code>	existing <code>collectPropsInAndChain</code> ; new <code>Proj.ProjChain</code> ; new lemmas <code>ne_symm_eq</code> , <code>distinct_bool_false</code>
<code>semilattice_simp</code> <code>boolean_group_simp</code>	new <code>AciNorm.proveEq</code> new <code>AciNorm.proveXorEq</code>
<code>assoc_simp</code> <code>evaluate</code>	existing <code>Meta.AC.rewriteUnnormalizedTop</code> existing <code>synthDecidableInstance</code> ; new <code>decideProofOfNative</code> , <code>proveByCases</code> , <code>decideGround</code>
<code>rare_rewrite</code>	dispatched to <code>Rare.lean</code> , below

5.6 RARE rewrites (`Rare.lean` 109, `Rare/Basic.lean` 44, `Rare/Rules.lean` 1 380 lines)

A `rare_rewrite` step names a rule of the RARE file Carcara was run with; the `core-simp-rare` pass replays every `*_simplify` step as a chain of such steps and `evaluate`. The generator `gen.py` reads that file, `rewrites.eo`, and emits one theorem per rule, two for rules polymorphic over `Int` and `Real`, plus a dispatch table describing each rule's arguments; existing proofs are preserved across regenerations. `Rare.lean` instantiates the theorem with the step's arguments by a reducing telescope and unification.

Leaving aside the two non-linear `mult-tangent` rules, the table holds 105 rules and 113 theorems. Their proofs are almost entirely existing lemmas:

Proof of the generated theorem	Theorems
existing <code>Int.Rewrite.*</code> and <code>Rat.Rewrite.*</code> lemmas (<code>Int/Rewrites</code> , <code>Rat/Rewrites</code>)	45
existing <code>Prop.bool_*</code> and <code>Prop.ite_*</code> lemmas (<code>Prop/Rewrites</code>)	41
existing <code>Builtin.ite_*</code> and <code>UF.eq_*</code> , <code>UF.distinct_binary_elim</code> (<code>Builtin/Rewrites</code> , <code>UF/Rewrites</code>)	13
<code>rfl</code>	8
Lean core <code>Int.emod_emod</code> , <code>Int.emod_def</code> , or an inline <code>propext</code>	3
short tactic proof	2
new lemma <code>or_not_refl</code> (a rule newer than <code>lean-cvc5</code> 's enumeration)	1

Six `array-*` rules have no statement, there being no array sort support, and are trusted. The rule `distinct-false` is special-cased in `UF.lean` through `Proj.mkAndProj`.

5.7 Quantifier rules (`Quant.lean`, 655 lines)

Rule	Component
<code>qnt_rm_unused</code> ($\forall$)	existing <code>Quant.reconstructRewrite</code> at <code>QUANT_UNUSED_VARS</code>
<code>qnt_join</code>	existing <code>Quant.reconstructRewrite</code> at <code>QUANT_MERGE_PRENEX</code>
<code>miniscope_distribute</code>	existing <code>Quant.reconstructRewrite</code> at <code>QUANT_MINISCOPE_AND</code>
<code>miniscope_split</code>	existing <code>Quant.reconstructRewrite</code> at <code>QUANT_MINISCOPE_OR</code>
<code>miniscope_ite</code>	existing <code>Quant.reconstructRewrite</code> at <code>QUANT_MINISCOPE_ITE</code>
<code>forall_inst</code>	existing <code>Prop.impliesElim</code> ; new <code>alignInstances</code>
<code>bind</code>	existing <code>Quant.exists_congr_eq</code> , Lean core <code>forall_congr</code> ; new lemma <code>orN_forall</code>
<code>onepoint</code>	new <code>Proj.mkOrInj</code> , <code>Proj.mkAndProj</code> ; existing <code>collectPropsInOrChain</code> , <code>collectPropsInAndChain</code> ; new lemmas <code>onepoint_forall</code> , <code>onepoint_exists</code>
<code>sko_ex</code> , <code>sko_forall</code>	new lemmas <code>sko_ex_eq</code> , <code>sko_forall_eq</code>
<code>qnt_simplify</code>	new lemmas <code>qnt_forall_true/false</code> , <code>qnt_exists_true/false</code>
<code>let</code>	new meta code, a <code>mkCongrArg</code> chain over anchor <code>lets</code>
<code>qnt_rm_unused</code> ($\exists$)	new meta code; the shared reconstruction is $\forall$ -only

The first five rows are handled with no new proof code at all: the Alethe rule name is mapped to the `cvc5` rewrite rule and the existing reconstructor is called through a `RewriteStep`. The anchor machinery (kept variables introduced first, assignments `let`-bound so that the substitution is definitional) is new and lives in the driver.

5.8 Assumptions (`Polyeq.lean`, 145 lines)

An `assume` must be an assertion of the problem. Carcara's `polyeq` pass makes that syntactic on its side: an `assume` matching an assertion only up to a flipped equality, a regrouped associative application or renamed bound variables is replaced by the assertion, with explicit steps deriving the proof's spelling. One difference survives the pass because Carcara's parser erases it: numerals are identified by value, so `0.0` and `0/1` are one constant to Carcara, the `assume` is an exact match there, nothing is elaborated, and the printer chooses its own spelling. The checker realizes both sides through `cvc5`'s parser, which is syntactic, and sees two terms. `Polyeq.lean` accepts exactly that: `polyeq` compares two `cvc5` terms structurally with numerals by value, and `polyeqProof?` proves the propositions equal by congruence down to the numerals, where the kernel unfolds one spelling into the other. A flipped equality or a renamed binder is rejected, as `Test/Alethe/M0/assume_flipped` checks. The path runs only after the exact and definitional matches fail; on the sample run in the repository no assumption reached it.

5.9 The new lemmas

`Alethe/Lemmas.lean` holds 52 theorems. They cover the clausal shapes of Alethe conclusions (`orN_of_impliesN`, `orN_of_impliesN_not`, `orN_of_map`, `orN_of_getElem`, `orN_forall`), the

Farkas, disequality, totality and `la_rw_eq` statements in `Int` and `Rat`, nine mixed-sort statements for `poly_simp_rel`, the Skolemization and one-point equalities, the four `qnt_simplify` constants, the three `connective_def` shapes, the `div_intro` family, and small facts about `distinct`. None duplicates a library lemma; most are thin wrappers over the existing `Int` and `Rat` lemmas named above.

5.10 What has no counterpart in the library

Genuinely Alethe-specific new work, with no `lean-smt` component to lean on: the Farkas driver of `la_generic`; the anchor and binder rules `bind`, `let`, `onpoint`, `ske_ex`, `ske_forall`; `distinct_elim`'s pair matching; the numeral matching of assumptions; the term-orientation logic of `trans`, `symm` and `cong` under anchor `lets`; and the step-by-step kernel driver of `Reconstruct.lean`, 775 lines, with its batching (`smt.alethe.batch`, `smt.alethe.batchSize`), concurrency (`smt.alethe.jobs`), inner re-check of a rejected batch, closed-term caching across scopes, and CSV timing output.

6 Tests, scripts and documents

`Test/Alethe/` adds 48 tests over 94 problem and proof files, grouped as `M0`, `QF_UF`, `UF`, `QF_LIA`, `QF_LRA`, `QF_UFLIA` and `Tactic`; each test is a single `##check_alethe` command over a problem and its proof, so the Lean side is 256 lines and the proof text 15037. `Test/Unit/AciNorm.lean` and `Test/Unit/ResNorm.lean` exercise the new normalizers directly; `Test/Unit/Normalize.expected` was refreshed because the `polyNorm` change moves an internal counter.

`scripts/setup-alethe.sh` builds `cvc5`, `Carcara` and the no-Mathlib `lean-smt`; `scripts/check-alethe.sh` runs the checker over a proof with a heap cap, phase timings and optional per-step CSV; `rule-boxplots.py` draws the per-rule plots. `docs-alethe/` holds the evaluation report `alethe-lean-smt.tex` with its figures and the two notes `rule-audit.md` and `elaboration-opportunities.md`.